



Outline

Overfitting

Regularizers

Introduction


- ▶ Training: learning model parameters.
- ▶ Validation: evaluate model on unseen data.
- ▶ Expect similar performance in training and validation data (both must come from the same underlying distribution).
- ▶ Chance hyper-parameters in the presence of performance gap.



Overfitting

“Model fits training data extremely good, but fails to generalize for unseen data”.

Common solution: give up some performance level on the training set in favor of improvement in the validation set.




Bias and variance

- ▶ **Bias:** Inability of a model to learn the true mapping relationship between x and y , which implies underfitting, e.g., straight line.
- ▶ **Variance:** Difference in fits between training and validation sets, which might imply overfitting, e.g., high-order polynomial.

Bias-variance trade-off

Ideal model: Low bias (accurate mapping function) and low variance (performance is consistent between training and validation sets).



Common practices

To reduce underfitting and overfitting:

underfitting

- ▶ Add capacity.
- ▶ Add epochs.
- ▶ Add data.
- ▶ Add features.
- ▶ Data augmentation.

overfitting

- ▶ Early stopping.
- ▶ Gather more data.
- ▶ Decrease features.
- ▶ Add regularizers.
- ▶ Data augmentation.

Outline

Overfitting

Regularizers

Introduction

Limit the model capacity, by adding a penalty to the parameters.

$$\mathcal{L}_T(x, y; \Omega) = \mathcal{L}(x, y; \Omega) + \alpha \mathcal{P}(\Omega),$$

where,

- ▶ $\mathcal{L}(x, y; \Omega)$ corresponds to the target loss already known,
- ▶ $\mathcal{P}(\Omega)$ is the penalty function on the parameters,
- ▶ α weights the penalty of the parameters, and
- ▶ $\mathcal{L}_T(x, y; \Omega)$ indicates the total loss.

Weight decay $L2$

$$y = w_1 x_1 + w_2 x_2 + \underline{w_3 x_1 x_2}$$

a.k.a., Ridge regression or Tikhonov regularization.

$$\mathcal{P}(\Omega) = \frac{1}{2} \|\Omega\|_2^2.$$

$$\min E(y, \hat{y}) + \frac{1}{2} \|\Omega\|_2^2$$

$$\frac{\partial E}{\partial w_i} =$$

Derivative:

Ω .

Lasso L_1

$$\mathcal{P}(\Omega) = \|\Omega\|_1 = \sum_i |\omega_i|.$$

Derivative:

$$\text{sign}(\Omega).$$

Besides keeping small values for ω_i , it also induces sparsity.

Elastic $L2L1$

$$\min \mathcal{E}(y, \hat{y}) + \alpha \mathcal{P}(\Omega)$$

Ridge + Lasso.


$$\mathcal{P}(\Omega) = \alpha_R \|\Omega\|_2^2 + \alpha_L \|\Omega\|_1.$$

Derivative:

$$\Omega + \text{sign}(\Omega).$$


Regularized space

Keep parameter values at the intersection between the Loss space and the Penalty space.




Dropout

- ▶ Add stochasticity.
- ▶ During training, limit the capacity of the model at random.
- ▶ Deactivate neurons (dropout) or weights (dropconnect).
- ▶ Forces the network to become redundant.
- ▶ Also helps make the model robust against variations.



(a) Standard Neural Network.



(b) Network applying dropout

Srivastava et al., 2014. “Dropout: A Simple Way to Prevent Neural Networks from Overfittin”.


Batch normalization, I

a.k.a., batchnorm.

- ▶ Normalization is often applied to input layer (accelerates learning).
- ▶ We can do the same to hidden layers.
- ▶ Adds noise and learns to be robust against it.
- ▶ Induces independence between layers.

Ioffe & Szegedy, 2015. ICML. “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift”.

Batch normalization, II



Compute the batch mean and variance:

$$\underline{\mu}_B = \frac{1}{M} \sum_{m=1}^M x^{(m)}, \quad \underline{\sigma}_B^2 = \frac{1}{M} \sum_{m=1}^M \left(x^{(m)} - \mu_B \right)^2.$$

Normalize and feed to next layer:

$$\hat{x}^{(m)} = \gamma \cdot \frac{x^{(m)} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} + \beta.$$

